


Alexandru Calotoiu

DaFIEx





# To\_GPU

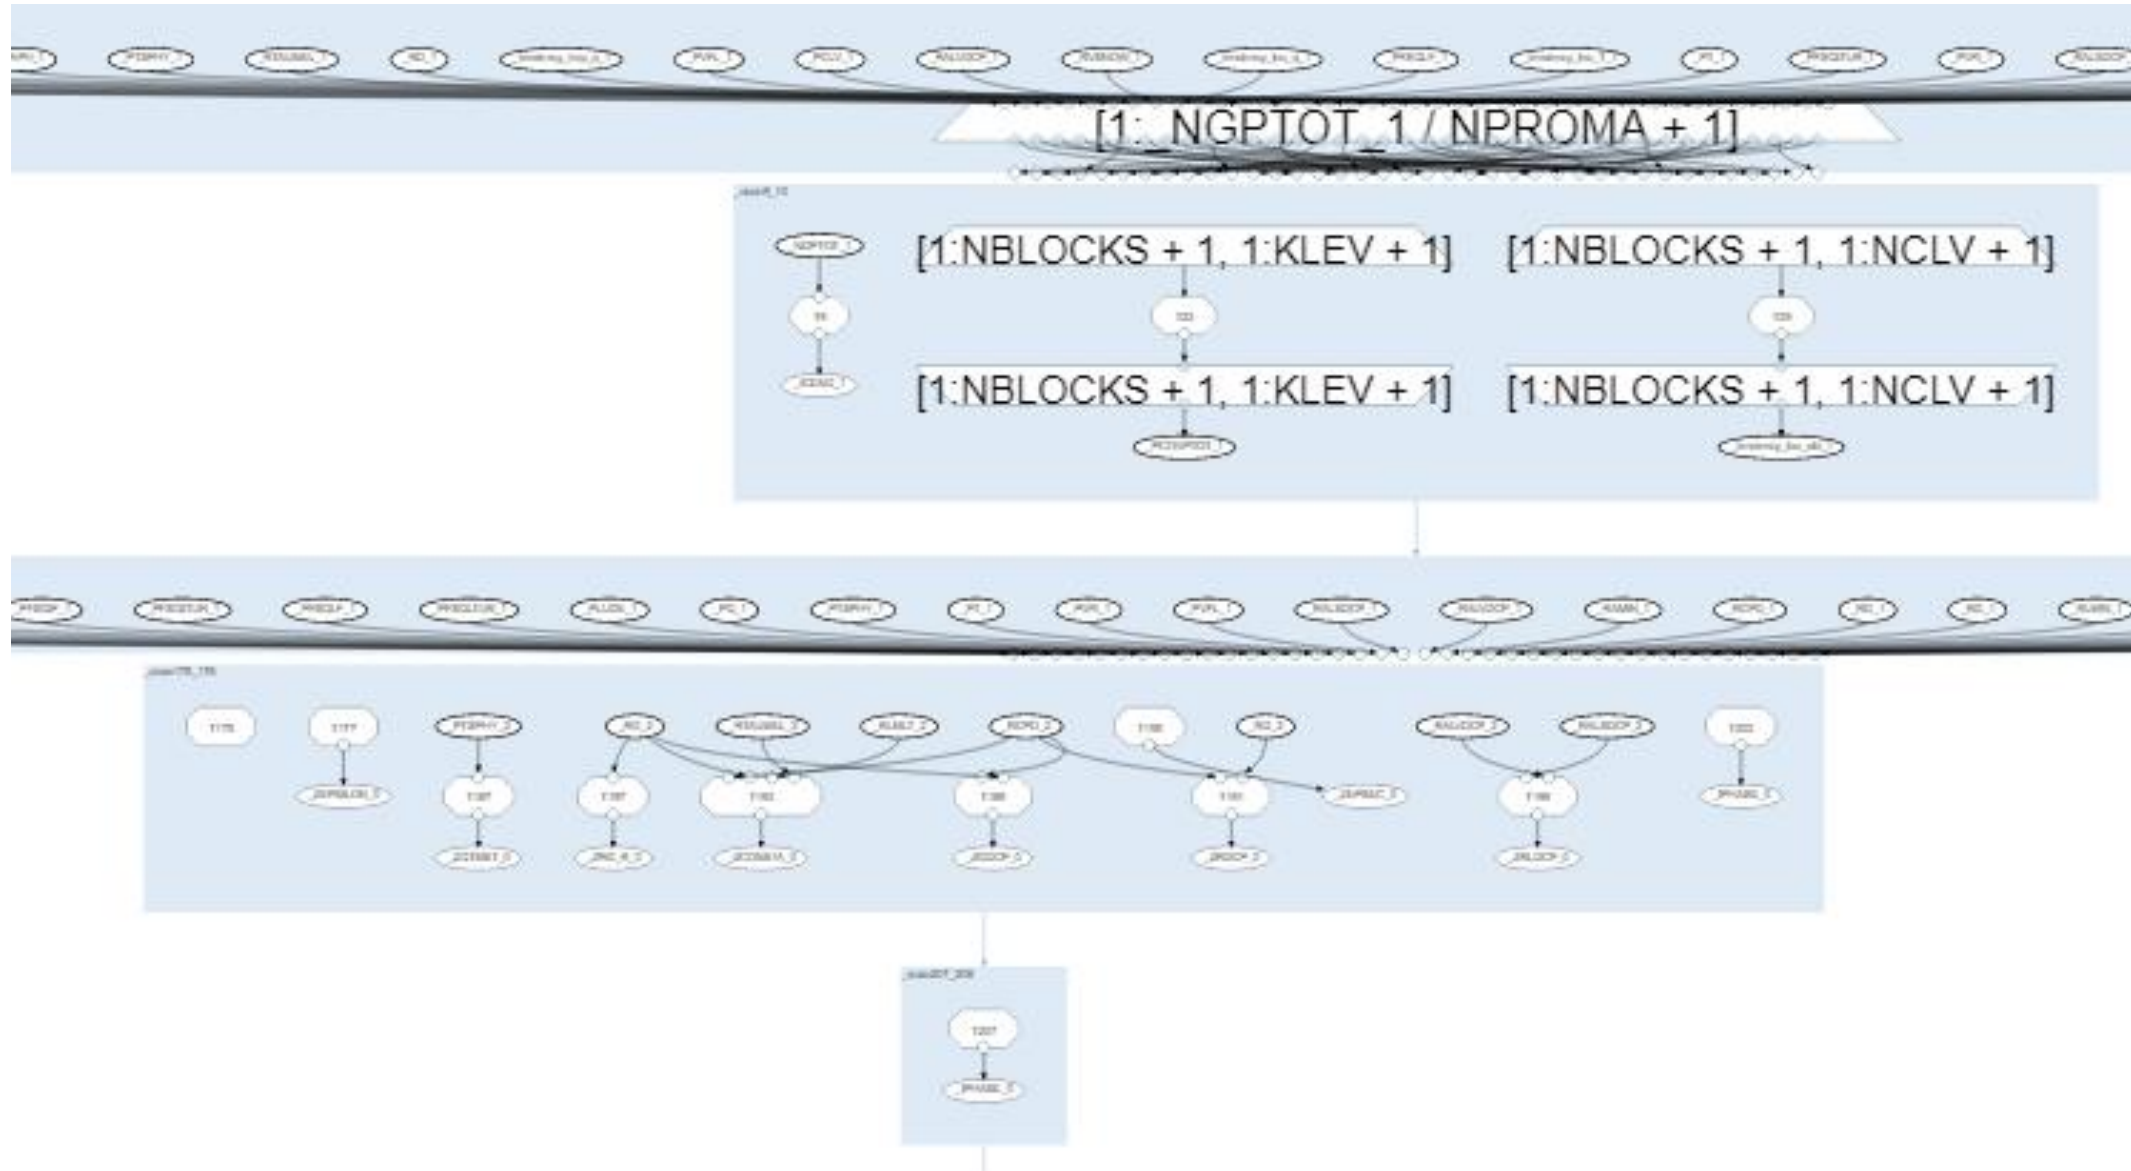
## Existing implementation outdated

- Tested on smaller kernels
- Works well on small kernels
- Not suited for large applications

## Research questions

- Which Maps to offload?
- What data to move?
- Should scalars or tasklets be moved?

# Example



# Correctness

## Improve to\_GPU to handle

- Nested SDFGs
- Scalars
- Single Tasklets

## Additional requirements

- Map Fission
- Heuristic Map Fusion