

Berke Ates, Prof. Dr. Torsten Hoefler, Dr. Tal Ben-Nun, Dr. Alexandru Calotoiu

MLIR-SDFG: A Data-Centric Dialect for MLIR



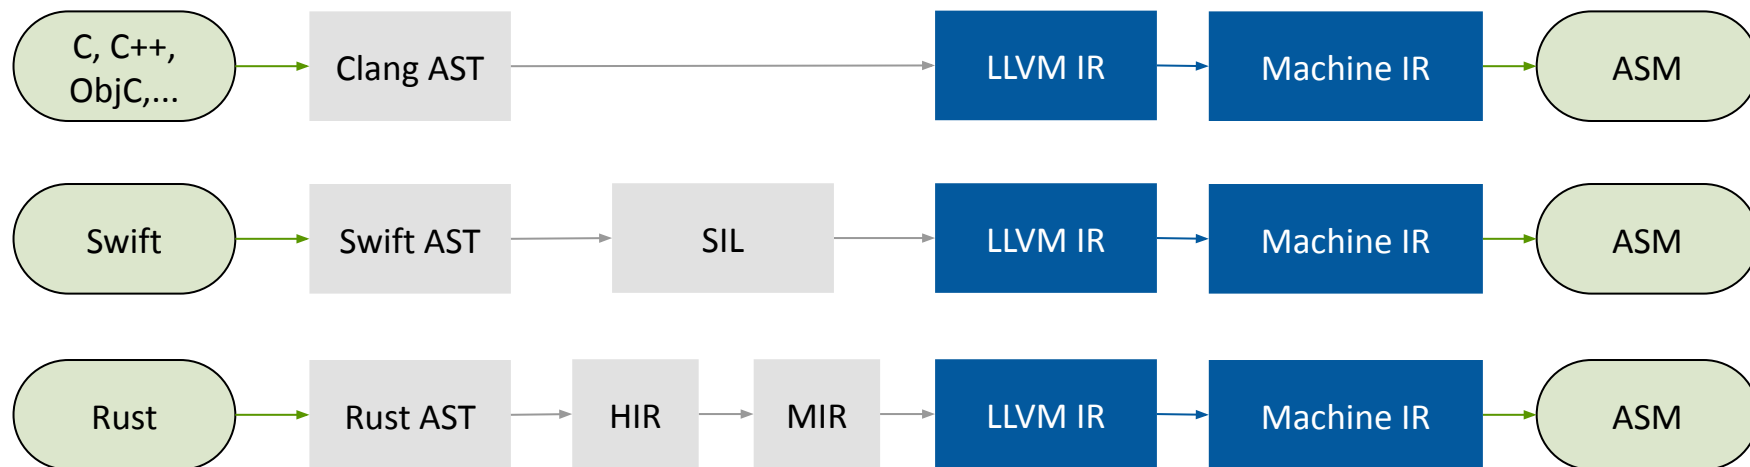


Motivation



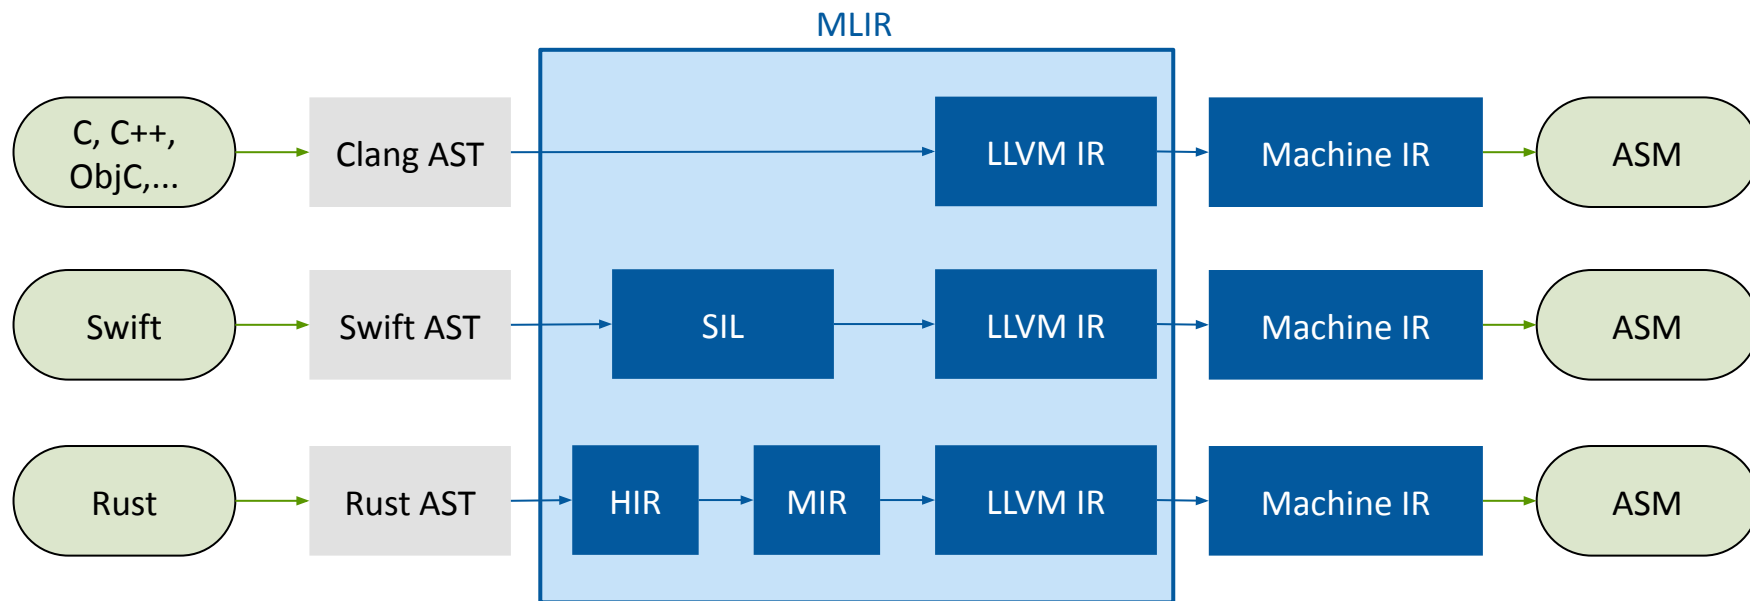


Motivation



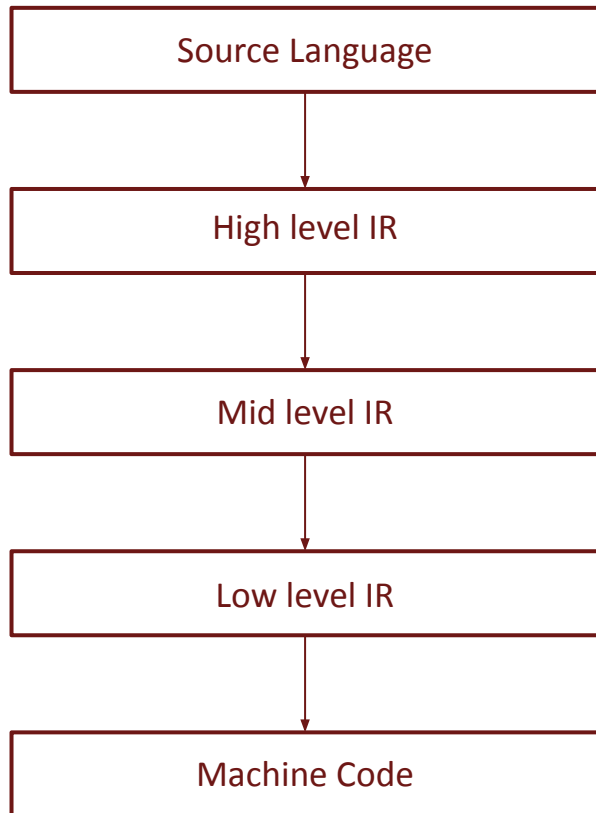


Motivation

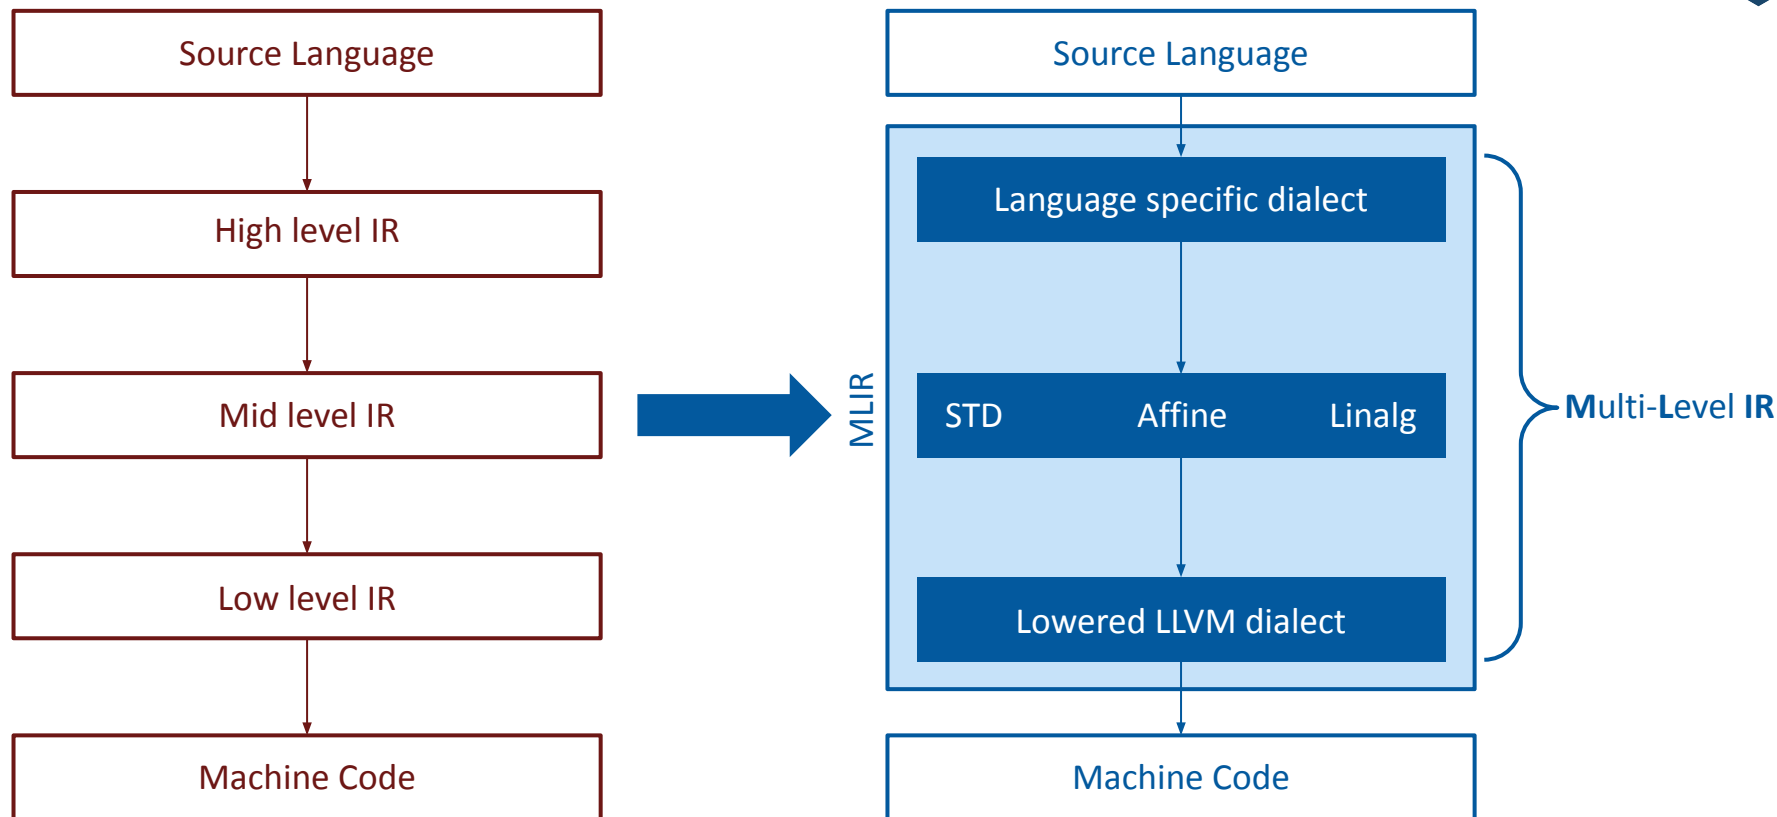




Compiling with MLIR



Compiling with MLIR





Dialects: Logical groups

- Define a namespace
- List of custom types
- List of operations
- Custom parser, printer and verifier
- List of passes: Analysis, Transformation, Dialect conversion

Dialects can coexist



```

affine.for %j = 0 to %nk {

    %ij = arith.muli %i, %j : index
    %ij_1 = arith.addi %ij, %1 : index
    %ij_1_64 = arith.index_cast %ij_1 : index to i64
    %ni_64 = arith.index_cast %ni : index to i64

    %rem = llvm.urem %ij_1_64, %ni_64 : i64
    %rem_f = llvm.bitcast %rem : i64 to f64
    %ni_64_f = llvm.bitcast %ni_64 : i64 to f64
    %entry = llvm.fdiv %rem_f, %ni_64_f : f64

    memref.store %entry, %A[%i, %j] : memref<?x?xf64>

}
    
```


Dialects can coexist



Affine affine.for %j = 0 to %nk {

```
%ij = arith.muli %i, %j : index
%ij_1 = arith.addi %ij, %1 : index
%ij_1_64 = arith.index_cast %ij_1 : index to i64
%ni_64 = arith.index_cast %ni : index to i64
```

```
%rem = llvm.urem %ij_1_64, %ni_64 : i64
%rem_f = llvm.bitcast %rem : i64 to f64
%ni_64_f = llvm.bitcast %ni_64 : i64 to f64
%entry = llvm.fdiv %rem_f, %ni_64_f : f64
```

```
memref.store %entry, %A[%i, %j] : memref<?x?xf64>
```

Affine }



Dialects can coexist

Affine

```
affine.for %j = 0 to %nk {
```

Arith

```
%ij = arith.muli %i, %j : index
%ij_1 = arith.addi %ij, %1 : index
%ij_1_64 = arith.index_cast %ij_1 : index to i64
%ni_64 = arith.index_cast %ni : index to i64
```

```
%rem = llvm.urem %ij_1_64, %ni_64 : i64
%rem_f = llvm.bitcast %rem : i64 to f64
%ni_64_f = llvm.bitcast %ni_64 : i64 to f64
%entry = llvm.fdiv %rem_f, %ni_64_f : f64
```

```
memref.store %entry, %A[%i, %j] : memref<?x?xf64>
```

Affine }

Dialects can coexist



Affine	affine.for %j = 0 to %nk {
Arith	%ij = arith.muli %i, %j : index %ij_1 = arith.addi %ij, %1 : index %ij_1_64 = arith.index_cast %ij_1 : index to i64 %ni_64 = arith.index_cast %ni : index to i64
LLVM	%rem = llvm.urem %ij_1_64, %ni_64 : i64 %rem_f = llvm.bitcast %rem : i64 to f64 %ni_64_f = llvm.bitcast %ni_64 : i64 to f64 %entry = llvm.fdiv %rem_f, %ni_64_f : f64
	memref.store %entry, %A[%i, %j] : memref<?x?xf64>
Affine	}



Dialects can coexist

- Code reuse
- Dialects can target specific problems
- Dialects can optimize for specific hardware

Affine

```
affine.for %j = 0 to %nk {
```

Arith

```
%ij = arith.muli %i, %j : index
%ij_1 = arith.addi %ij, %1 : index
%ij_1_64 = arith.index_cast %ij_1 : index to i64
%ni_64 = arith.index_cast %ni : index to i64
```

LLVM

```
%rem = llvm.urem %ij_1_64, %ni_64 : i64
%rem_f = llvm.bitcast %rem : i64 to f64
%ni_64_f = llvm.bitcast %ni_64 : i64 to f64
%entry = llvm.fdiv %rem_f, %ni_64_f : f64
```

Memref

```
memref.store %entry, %A[%i, %j] : memref<?x?xf64>
```

Affine

```
}
```



Operations, Not Instructions

- No fixed instruction set
- Operations are “opaque functions”

```
%res = mydialect.opname {someAttribute = true} (%arg1) : i32
```



Operations, Not Instructions

- No fixed instruction set
- Operations are “opaque functions”

```
%res = mydialect.opname {someAttribute = true} (%arg1) : i32
```

↑
Name of
the result



Operations, Not Instructions

- No fixed instruction set
- Operations are “opaque functions”

Dialect prefix

↓

%res = mydialect.opname {someAttribute = true} (%arg1) : i32

↑

Name of
the result



Operations, Not Instructions

- No fixed instruction set
- Operations are “opaque functions”

Diagram illustrating an operation syntax:

```
%res = mydialect.opname {someAttribute = true} (%arg1) : i32
```

Annotations:

- Dialect prefix** (green arrow) points to `mydialect`.
- Op ID** (blue arrow) points to `opname`.
- Name of the result** (black arrow) points to `%res`.



Operations, Not Instructions

- No fixed instruction set
- Operations are “opaque functions”

The diagram illustrates the syntax of an MLIR operation: `%res = mydialect.opname {someAttribute = true} (%arg1) : i32`. The components are color-coded and labeled with arrows:

- mydialect** (green): Labeled "Dialect prefix" with a green arrow pointing to it.
- opname** (blue): Labeled "Op ID" with a blue arrow pointing to it.
- {someAttribute = true}** (purple): Labeled "List of attributes" with a purple arrow pointing to it.
- %res** (black): Labeled "Name of the result" with a black arrow pointing to it.



Operations, Not Instructions

- No fixed instruction set
- Operations are “opaque functions”

`%res = mydialect.opname {someAttribute = true} (%arg1) : i32`

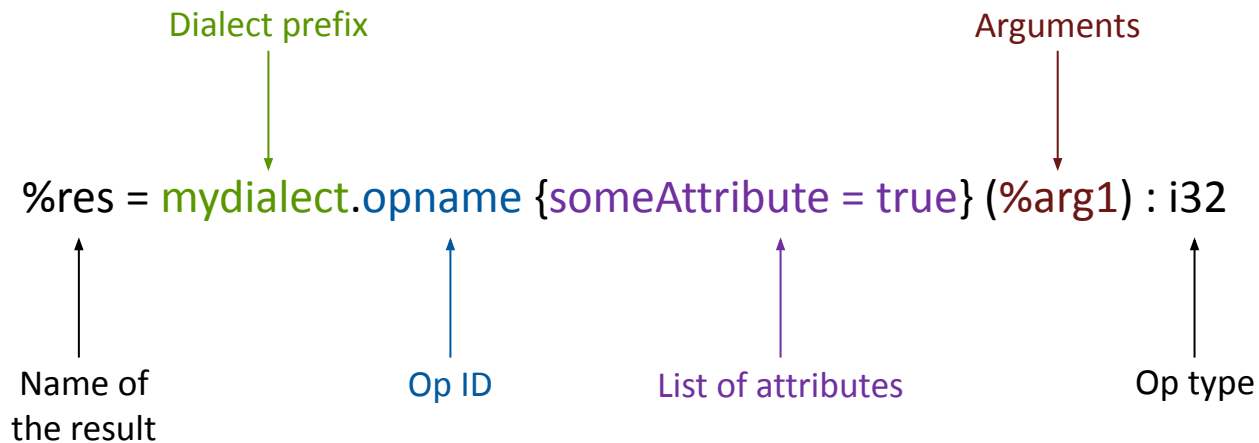
Diagram illustrating the components of an operation in SPCL:

- Dialect prefix** (green arrow) points to `mydialect`.
- Op ID** (blue arrow) points to `opname`.
- List of attributes** (purple arrow) points to `{someAttribute = true}`.
- Arguments** (red arrow) points to `(%arg1)`.
- Name of the result** (black arrow) points to `%res`.



Operations, Not Instructions

- No fixed instruction set
- Operations are “opaque functions”



Recursive nesting: Operations -> Regions -> Blocks



```
mydialect.whileNZ (%i) ({  
    ^block1:  
        %res = mydialect.add %a, %b : i32  
        br ^block2  
  
    ^block2:  
        // Some more code  
  
})
```

Recursive nesting: Operations -> Regions -> Blocks



```
Op mydialect.whileNZ (%i) ({  
    ^block1:  
        %res = mydialect.add %a, %b : i32  
        br ^block2  
  
    ^block2:  
        // Some more code  
  
})
```



Recursive nesting: Operations -> Regions -> Blocks

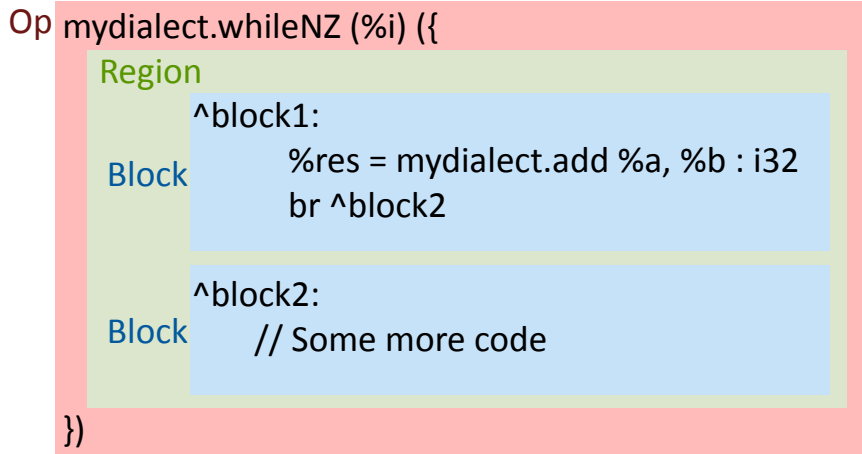
- Operations may contain regions

```
Op mydialect.whileNZ (%i) ({  
  Region  
    ^block1:  
      %res = mydialect.add %a, %b : i32  
      br ^block2  
  
    ^block2:  
      // Some more code  
})
```

Recursive nesting: Operations -> Regions -> Blocks



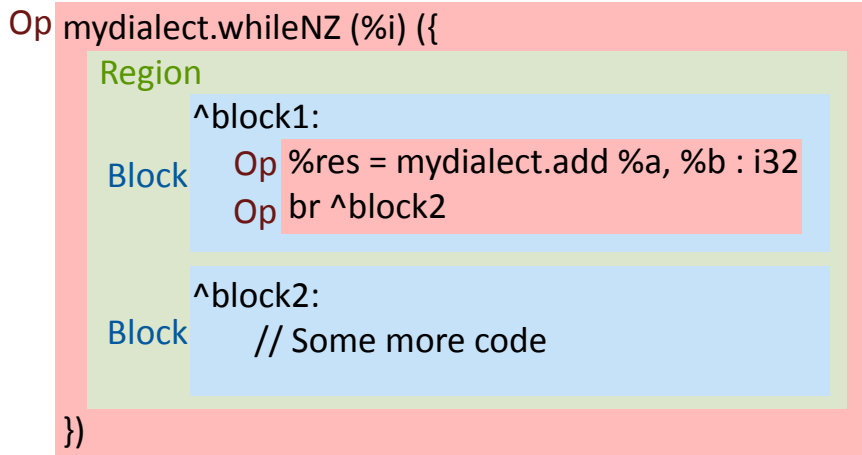
- Operations may contain regions
- Regions contain list of blocks



Recursive nesting: Operations -> Regions -> Blocks



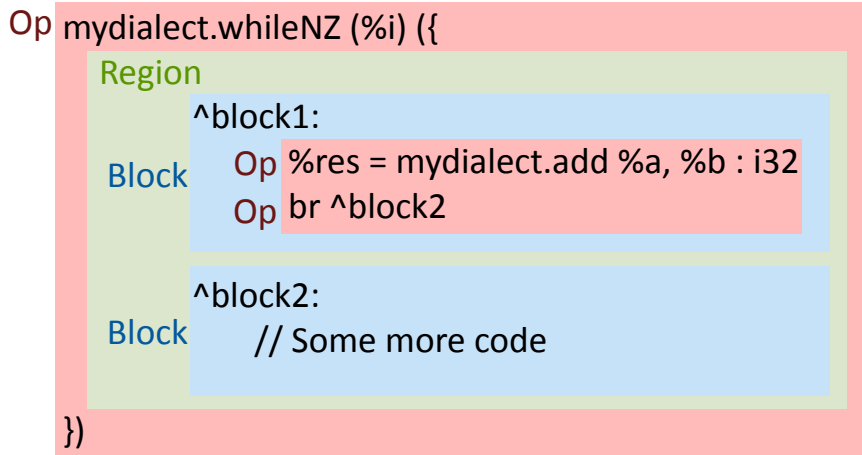
- Operations may contain regions
- Regions contain list of blocks
- Blocks contain list of operations



Recursive nesting: Operations -> Regions -> Blocks



- Operations may contain regions
- Regions contain list of blocks
- Blocks contain list of operations
- Allows modelling hierarchical structures



Recursive nesting: Operations -> Regions -> Blocks




- Operations may contain regions
- Regions contain list of blocks
- Blocks contain list of operations
- Allows modelling hierarchical structures




Everything is an operation


DaCe: **D**ata-**C**entric




DaCe: Data-Centric



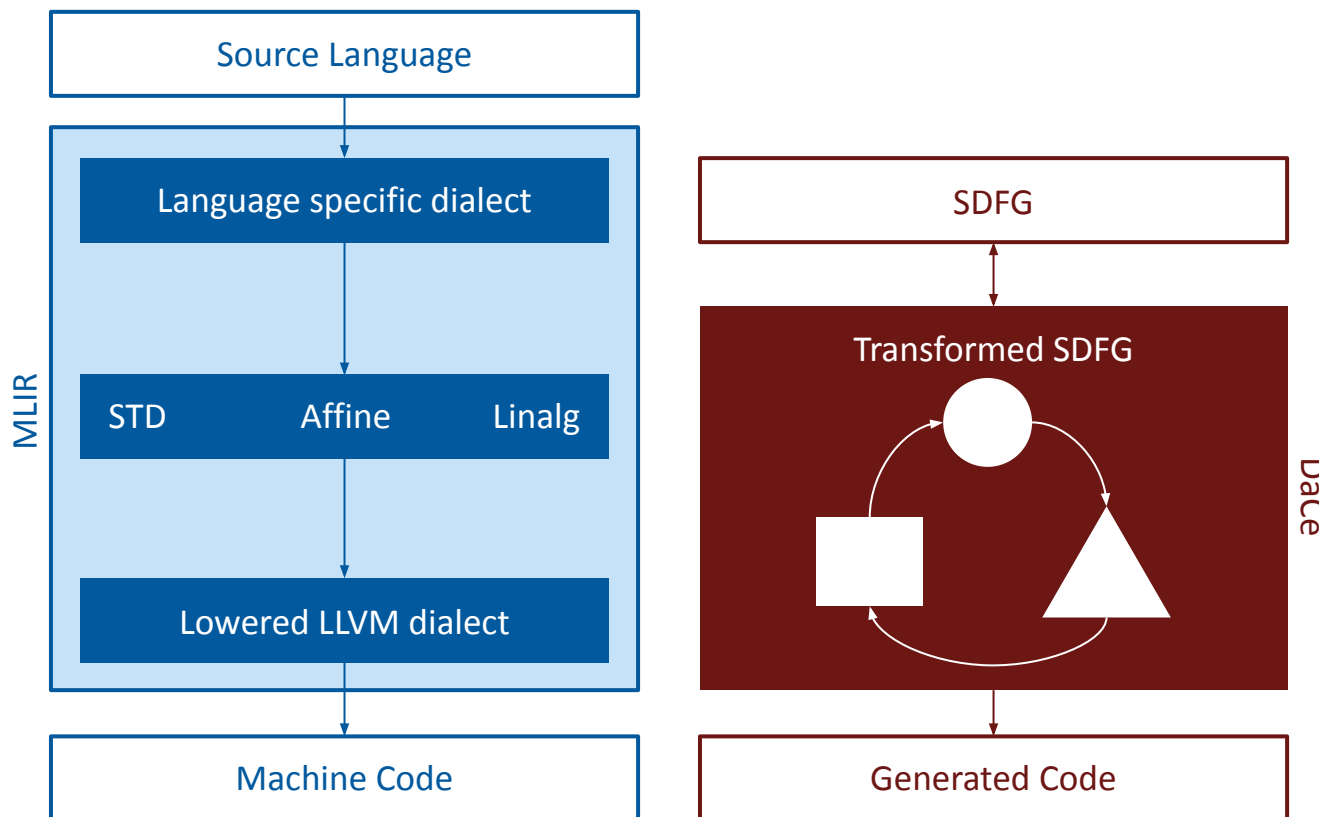
DaCe: Data-Centric



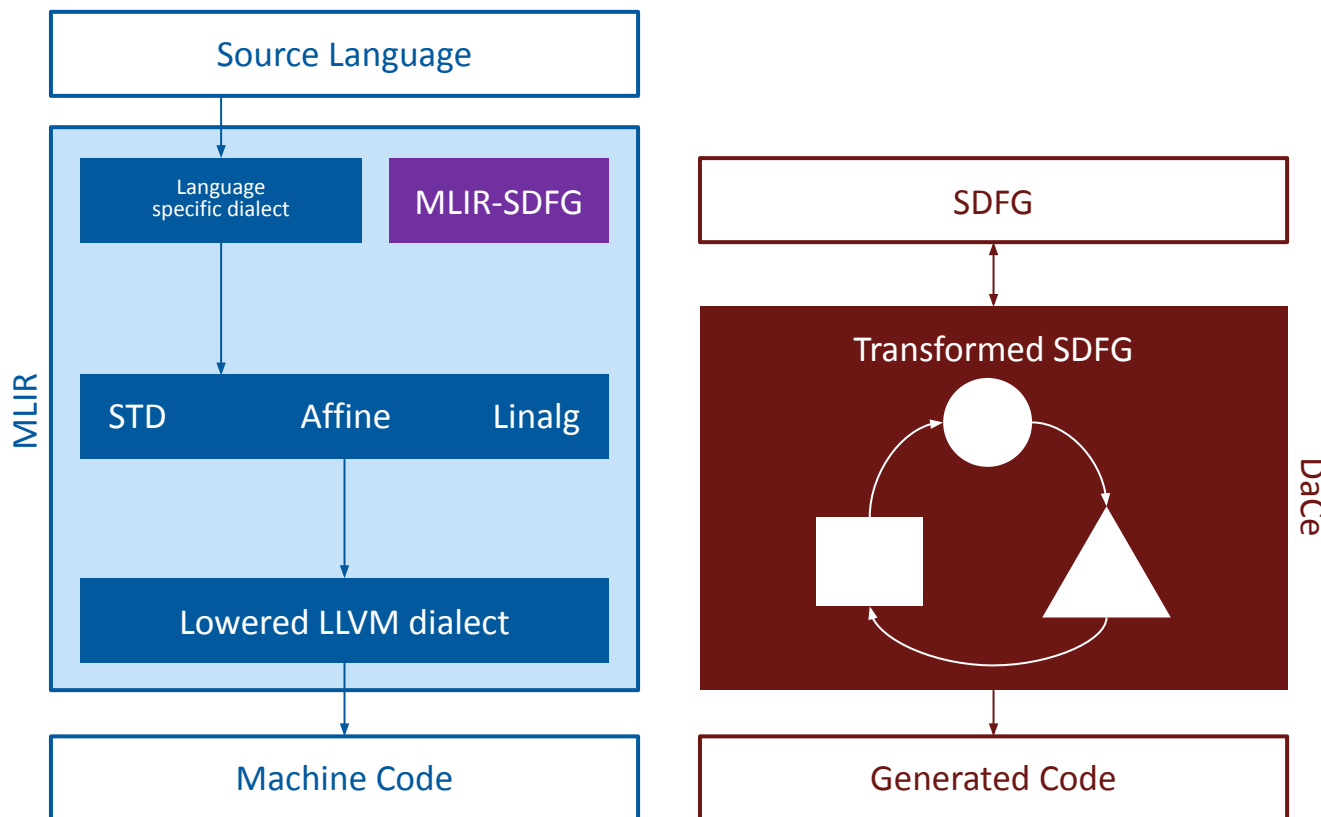
DaCe: Data-Centric



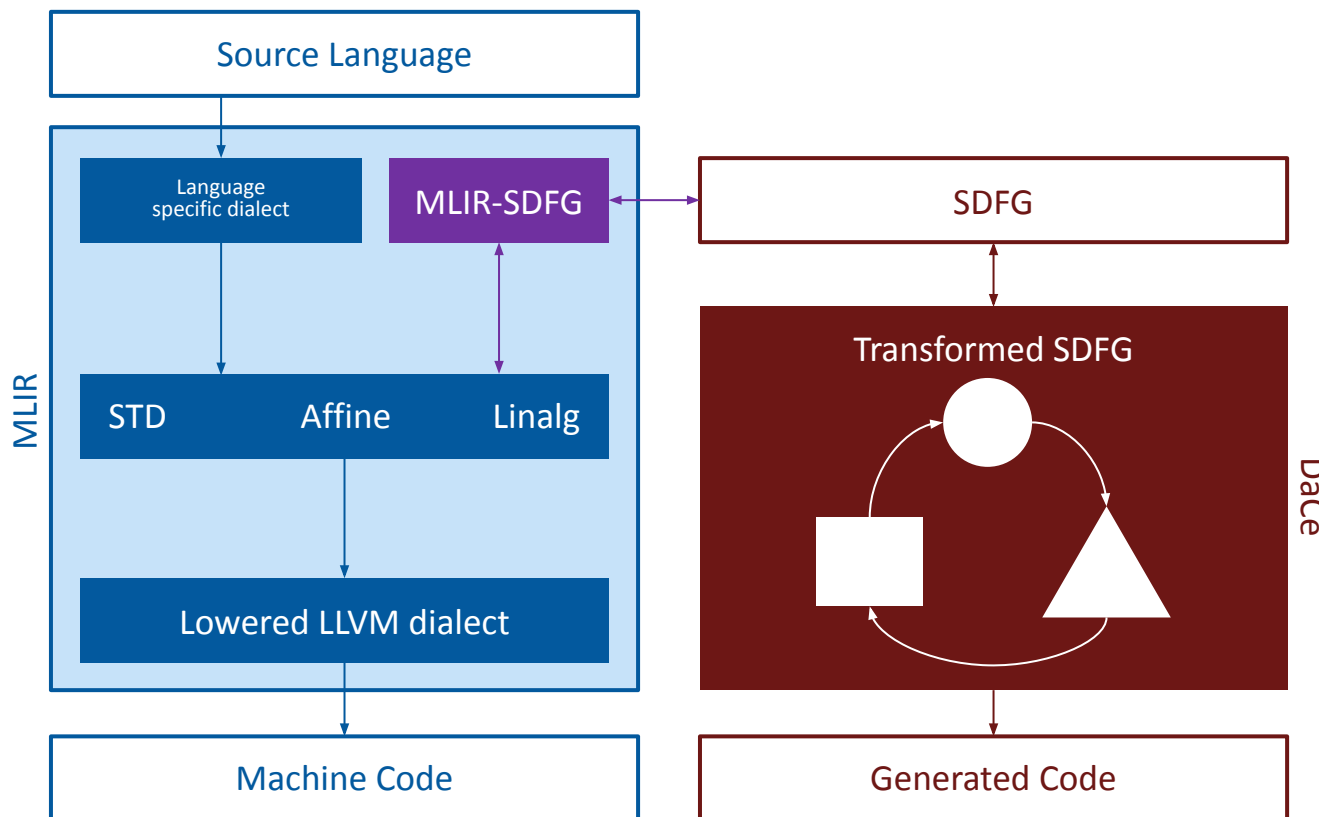
MLIR-SDFG: Between MLIR and DaCe



MLIR-SDFG: Between MLIR and DaCe



MLIR-SDFG: Between MLIR and DaCe



SDFG: Stateful DataFlow multiGraph




```
%A = sdfg.alloc() : !sdfg.array<2xi32>
```

```
%B = sdfg.alloc() : !sdfg.array<2xi32>
```

```
%a = sdfg.load %A[0]
```

```
sdfg.store %c, %B[0]
```



SDFG: Stateful DataFlow multiGraph




```
%A = sdfg.alloc() : !sdfg.array<2xi32>
%B = sdfg.alloc() : !sdfg.array<2xi32>
```


```
%a = sdfg.load %A[0]
```

```
%c = sdfg.tasklet @name(%a: i32) -> i32 {
  %r = // Some computation
  sdfg.return %r
}
```



```
sdfg.store %c, %B[0]
```



SDFG: Stateful DataFlow multiGraph



SDFG: Stateful DataFlow multiGraph




SDFG: Stateful DataFlow multiGraph

```
%A = sdfg.alloc() : !sdfg.array<sym("N")xi32>
%B = sdfg.alloc() : !sdfg.array<sym("N")xi32>
```

```
sdfg.map (%i) = (0) to (sym("N")) step (1) {
  %a = sdfg.load %A[%i]
```

```
  %c = sdfg.tasklet @name(%a: i32) -> i32 {
    %r = // Some computation
    sdfg.return %r
  }
```

```
sdfg.store %c, %B[%i]
}
```



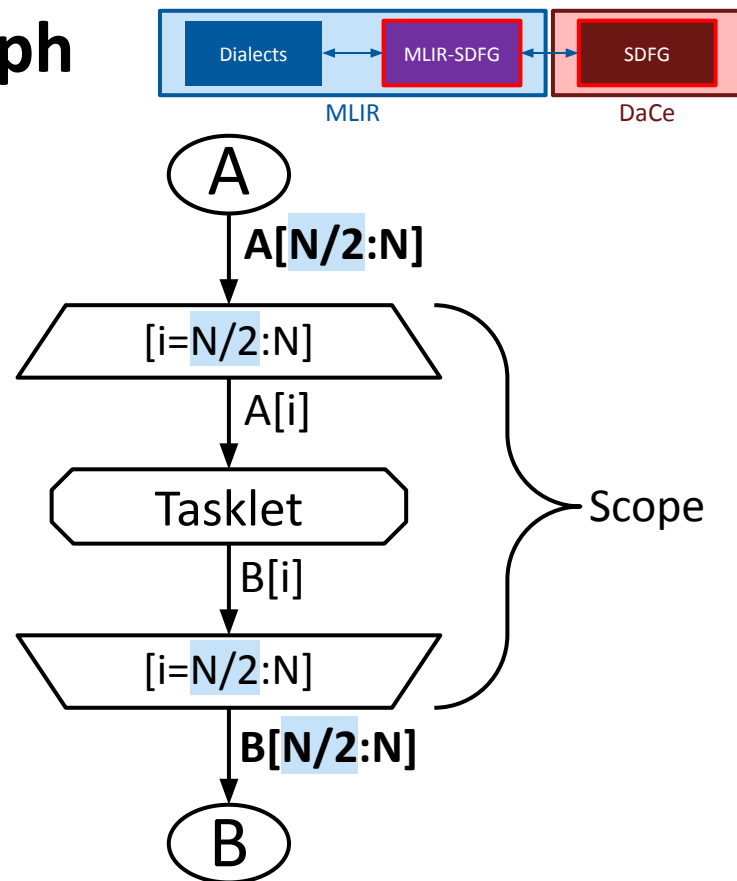
SDFG: Stateful DataFlow multiGraph

```
%A = sdfg.alloc() : !sdfg.array<sym("N")xi32>
%B = sdfg.alloc() : !sdfg.array<sym("N")xi32>



sdfg.map (%i) = (sym("N/2")) to (sym("N")) step (1) {
  %a = sdfg.load %A[%i]

  %c = sdfg.tasklet @name(%a: i32) -> i32 {
    %r = // Some computation
    sdfg.return %r
  }


  sdfg.store %c, %B[%i]
}
```



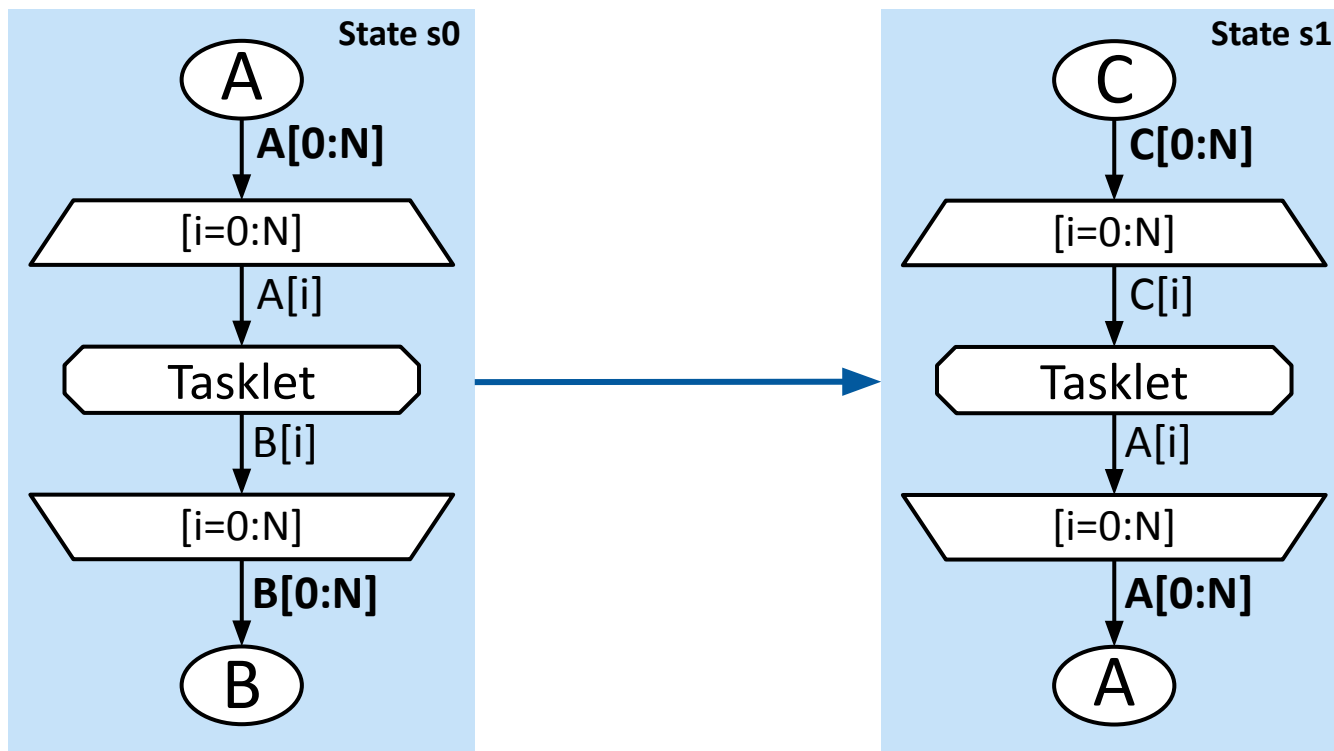
SDFG: Stateful DataFlow multiGraph



SDFG: Stateful DataFlow multiGraph



SDFG: Stateful DataFlow multiGraph



SDFG: Stateful DataFlow multiGraph




```
%C = sdfg.alloc() : !sdfg.array<sym("N")xi32>
%A = sdfg.alloc() : !sdfg.array<sym("N")xi32>
```

```
sdfg.state @s1 {
  sdfg.map (%i) = (0) to (sym("N")) step (1) {
    %c = sdfg.load %C[%i]

    %res = sdfg.tasklet @name(%c: i32) -> i32 {
      %r = // Some computation
      sdfg.return %r
    }

    sdfg.store %res, %A[%i]
  }
}
```




SDFG: Stateful DataFlow multiGraph

```
sdfg.sdfg{entry=@s0} {
  sdfg.state @s0 {...}
  sdfg.state @s1 {...}
  sdfg.state @s2 {...}
  sdfg.state @s3 {...}

```

```
  sdfg.edge{assign=["i: 0"]} @s0 -> @s1
  sdfg.edge{condition="i < 10"} @s1 -> @s2
  sdfg.edge{condition="i >= 10"} @s1 -> @s3
}
```




SDFG: Stateful DataFlow multiGraph

```

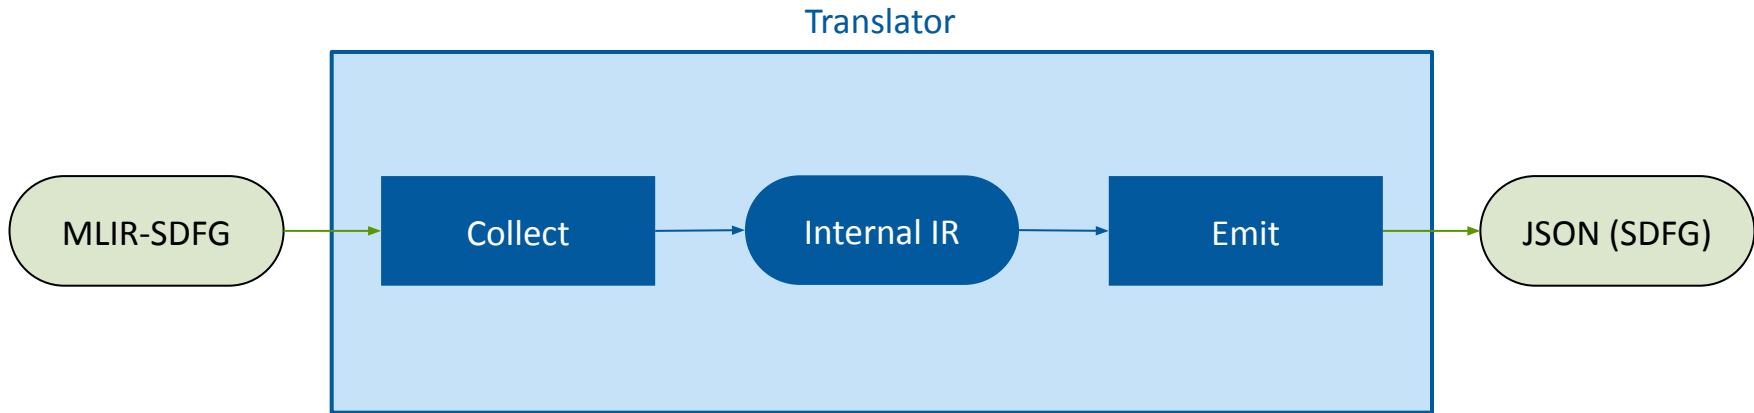
sdfg.sdfg{entry=@s0} {
  sdfg.state @s0 {...}
  sdfg.state @s1 {...}
  sdfg.state @s2 {...}
  sdfg.state @s3 {...}

  sdfg.edge{assign=["i: 0"]} @s0 -> @s1
  sdfg.edge{condition="i < 10"} @s1 -> @s2
  sdfg.edge{condition="i >= 10"} @s1 -> @s3
  sdfg.edge{assign=["i: i + 1"]} @s2 -> @s1
}

```



MLIR-SDFG -> SDFG Translator



Lifting to Python



```
sdfg.tasklet @add(%a, %b) -> f64 {  
    %r = arith.addi %a, %b  
    sdfg.return %r  
}
```

```
sdfg.tasklet @get_0() -> f64 {  
    %r = arith.constant 0.0  
    sdfg.return %r  
}
```



Lifting to Python



```
sdfg.tasklet @add(%a, %b) -> f64 {  
    %r = arith.addi %a, %b  
    sdfg.return %r  
}
```



__out = (a + b)

```
sdfg.tasklet @get_0() -> f64 {  
    %r = arith.constant 0.0  
    sdfg.return %r  
}
```



__out = dace.float64(0.0)



Optimizing Passes

```
sdfg.sdfg @name(%arg0: i32, %arg1: i32) {  
  ...  
  %c = sdfg.tasklet @add(%arg0, %arg1) -> (i32) {...}  
  ...  
}
```



```
sdfg.sdfg @name(%arg0: i32, %arg1: i32) {  
  %0 = sdfg.alloc {name = "_arg0"}() : !sdfg.array<i32>  
  %1 = sdfg.alloc {name = "_arg1"}() : !sdfg.array<i32>  
  ...  
  %a = sdfg.load %0[] : !sdfg.array<i32> -> i32  
  %b = sdfg.load %1[] : !sdfg.array<i32> -> i32  
  %c = sdfg.tasklet @add(%a, %b) -> (i32) {...}  
  ...  
}
```



Optimizing Passes

```
sdfg.sdfg @name(%arg0: i32, %arg1: i32) {
  ...
  %c = sdfg.tasklet @add(%arg0, %arg1) -> (i32) {...}
  ...
}
```



```
sdfg.sdfg @name(%arg0: i32, %arg1: i32) {
  %0 = sdfg.alloc {name = "_arg0"}() : !sdfg.array<i32>
  %1 = sdfg.alloc {name = "_arg1"}() : !sdfg.array<i32>
  ...
  %a = sdfg.load %0[] : !sdfg.array<i32> -> i32
  %b = sdfg.load %1[] : !sdfg.array<i32> -> i32
  %c = sdfg.tasklet @add(%a, %b) -> (i32) {...}
  ...
}
```



```
sdfg.state @state_1 {
  ...
  %8 = sdfg.load %5[] : !sdfg.array<index> -> index
}
```




```
sdfg.state @state_1 {
  ...
  // Removed dead load operation
}
```

Polygeist




Converter



MLIR-SDFG

Converter



Converter: Funcs



```
sdfg.sdfg @binops(%arg0: i32) {
```

```
func @binops(%arg0: i32) {  
  %c0 = arith.addi %arg0, %arg0 : i32  
  %c1 = arith.muli %c0, %arg0 : i32  
  return  
}
```




```
%c0 = arith.addi %arg0, %arg0 : i32
```

```
%c1 = arith.muli %c0, %arg0 : i32
```

```
}
```

Converter: Arith Operations



```
sdfg.sdfg @binops(%arg0: i32) {
```

```

  %0 = sdfg.tasklet @task_0(%arg0: i32) -> i32 {
    %c0 = arith.addi %arg0, %arg0 : i32
    sdfg.return %c0 : i32
  }

```


```

  %1 = sdfg.tasklet @task_1(%0: i32, %arg0: i32) -> i32 {
    %c1 = arith.muli %0, %arg0 : i32
    sdfg.return %c1 : i32
  }

```

```
}
```

Converter: Arith Operations



```
sdfg.sdfg @binops(%arg0: i32) {
```

```
  sdfg.state @state_0 {
```

```
    %0 = sdfg.tasklet @task_0(%arg0: i32) -> i32 {
```

```
      %c0 = arith.addi %arg0, %arg0 : i32
```

```
      sdfg.return %c0 : i32
```

```
    }
```

```
  }
```

```
  sdfg.state @state_1 {
```

```
    %1 = sdfg.tasklet @task_1(%0: i32, %arg0: i32) -> i32 {
```

```
      %c1 = arith.muli %0, %arg0 : i32
```



```
      sdfg.return %c1 : i32
```

```
    }
```

```
  }
```

```
}
```


Converter: Arith Operations



```

sdfg.sdfg @binops(%arg0: i32) {
  %tmp = sdfg.alloc {transient}() : !sdfg.array<i32>

```

```

sdfg.state @state_0 {
  %0 = sdfg.tasklet @task_0(%arg0: i32) -> i32 {
    %c0 = arith.addi %arg0, %arg0 : i32
    sdfg.return %c0 : i32
  }
}

```

```

sdfg.state @state_1 {
  %1 = sdfg.tasklet @task_1(%0: i32, %arg0: i32) -> i32 {
    %c1 = arith.muli %0, %arg0 : i32
    sdfg.return %c1 : i32
  }
}


```

```

}

```

Converter: Arith Operations



```

sdfg.sdfg @binops(%arg0: i32) {
  %tmp = sdfg.alloc {transient}(): !sdfg.array<i32>

```

```

sdfg.state @state_0 {
  %0 = sdfg.tasklet @task_0(%arg0: i32) -> i32 {
    %c0 = arith.addi %arg0, %arg0 : i32
    sdfg.return %c0 : i32
  }
}

```

```

sdfg.state @state_1 {
  %1 = sdfg.tasklet @task_1(%0: i32, %arg0: i32) -> i32 {
    %c1 = arith.muli %0, %arg0 : i32
    sdfg.return %c1 : i32
  }
}

```

```

sdfg.edge @state_0 -> @state_1
}

```

Converter: Memref Operations



```
%a = memref.load %A[%0] : memref<900xi32>
```



```
%a = sdfg.load %A[%0] : !sdfg.array<900xi32> -> i32
```

```
memref.store %a, %A[%0] : memref<900xi32>
```



```
sdfg.store %a, %A[%0] : i32 -> !sdfg.array<900xi32>
```

Converter: Memref Operations



```
%a = memref.load %A[%0] : memref<900xi32>
```



```
%a = sdfg.load %A[%0] : !sdfg.array<900xi32> -> i32
```

```
memref.store %a, %A[%0] : memref<900xi32>
```



```
sdfg.store %a, %A[%0] : i32 -> !sdfg.array<900xi32>
```

Converter: Memref Operations



```
%a = memref.load %A[%0] : memref<900xi32>
```



```
%a = sdfg.load %A[%0] : !sdfg.array<900xi32> -> i32
```


```
memref<?x900xi32>
```

```
memref.store %a, %A[%0] : memref<900xi32>
```



```
sdfg.store %a, %A[%0] : i32 -> !sdfg.array<900xi32>
```

Converter: Memref Operations



```
%a = memref.load %A[%0] : memref<900xi32>
```



```
%a = sdfg.load %A[%0] : !sdfg.array<900xi32> -> i32
```

```
memref.store %a, %A[%0] : memref<900xi32>
```



```
sdfg.store %a, %A[%0] : i32 -> !sdfg.array<900xi32>
```

```
memref<?x900xi32>
```



```
!sdfg.array<sym("N")x900xi32>
```

Converter: SCF For Loops

```
scf.for %arg11 = %c0 to %c1 step %c2 {  
  // Loop Body  
}
```




Converter: SCF For Loops

```
sdfg.state @init {}
sdfg.state @guard {}
sdfg.state @body {}
```

```
// Loop Body
```

```
sdfg.state @return {}
sdfg.state @exit {}
```

```
sdfg.edge{assign=["i: 0"]} @init -> @guard
sdfg.edge{condition="i < 10"} @guard -> @body
sdfg.edge{condition="i >= 10"} @guard -> @exit
sdfg.edge{assign=["i: i + 1"]} @return -> @guard
```



Experimental Setup


- PolyBench's 2mm, adi, gemver, heat-3d, trmm

Experimental Setup


- PolyBench's 2mm, adi, gemver, heat-3d, trmm
- CSCS Ault server
- Build compilation pipeline for MLIR-SDFG
- Compare with clang, gcc, polly, pluto
- Single- & Multi-Threaded
- Run 100 times on the large dataset
- 10 Runs warm-up

Architecture	x86_64
CPUs	72
Cores per socket	18
Base frequency	3.00 GHz
L1I cache size	32kB
L1D cache size	32kB
LLVM/MLIR version	14.0.0git
Clang version	10.0.1
DaCe version	0.11.4


Compilation Pipeline




Compilation Pipeline



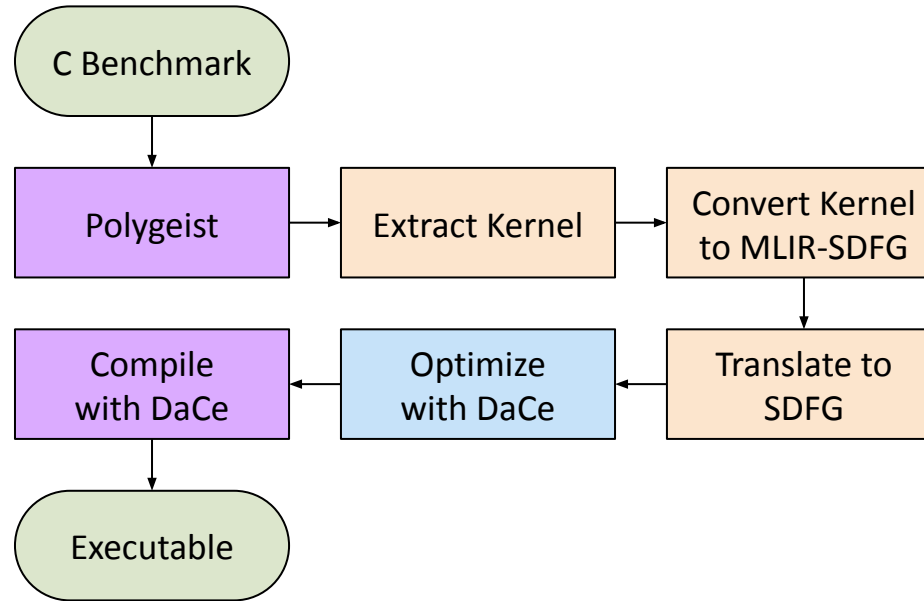
Compilation Pipeline



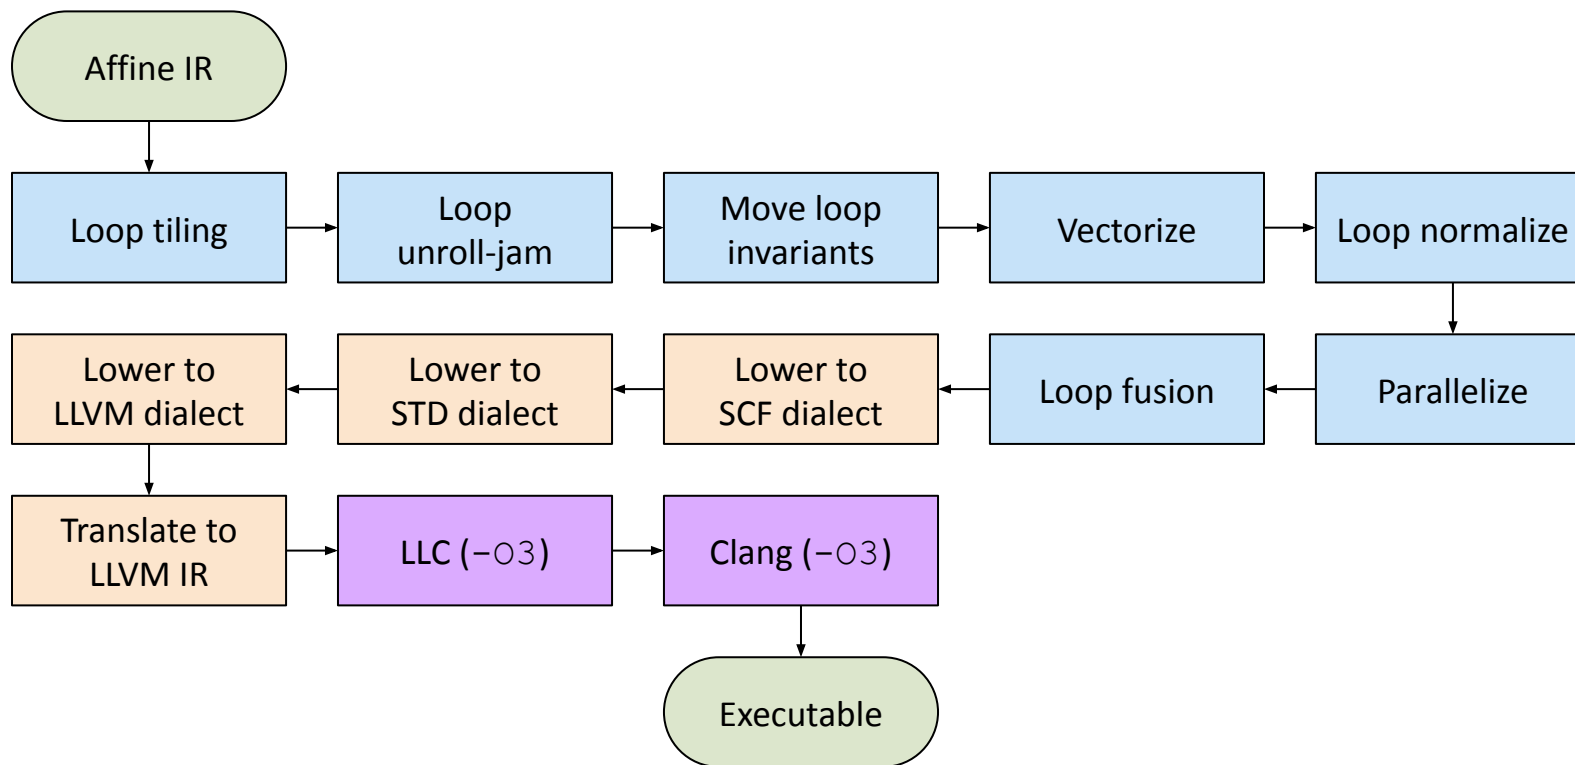
Compilation Pipeline




Compilation Pipeline



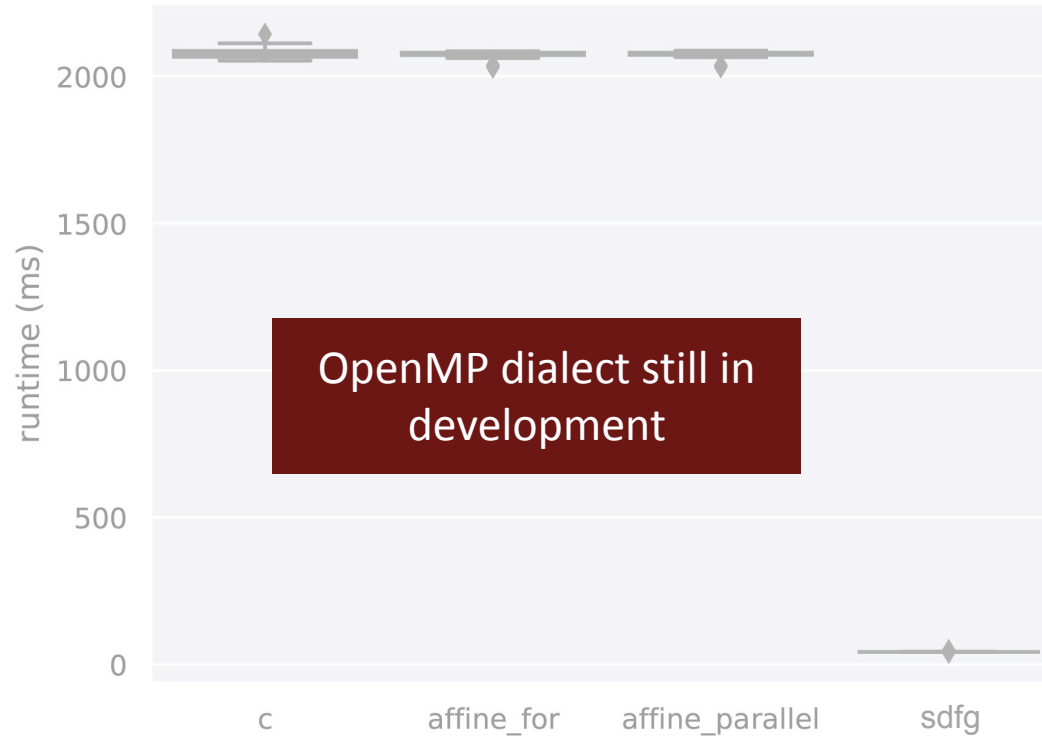
Affine pipeline



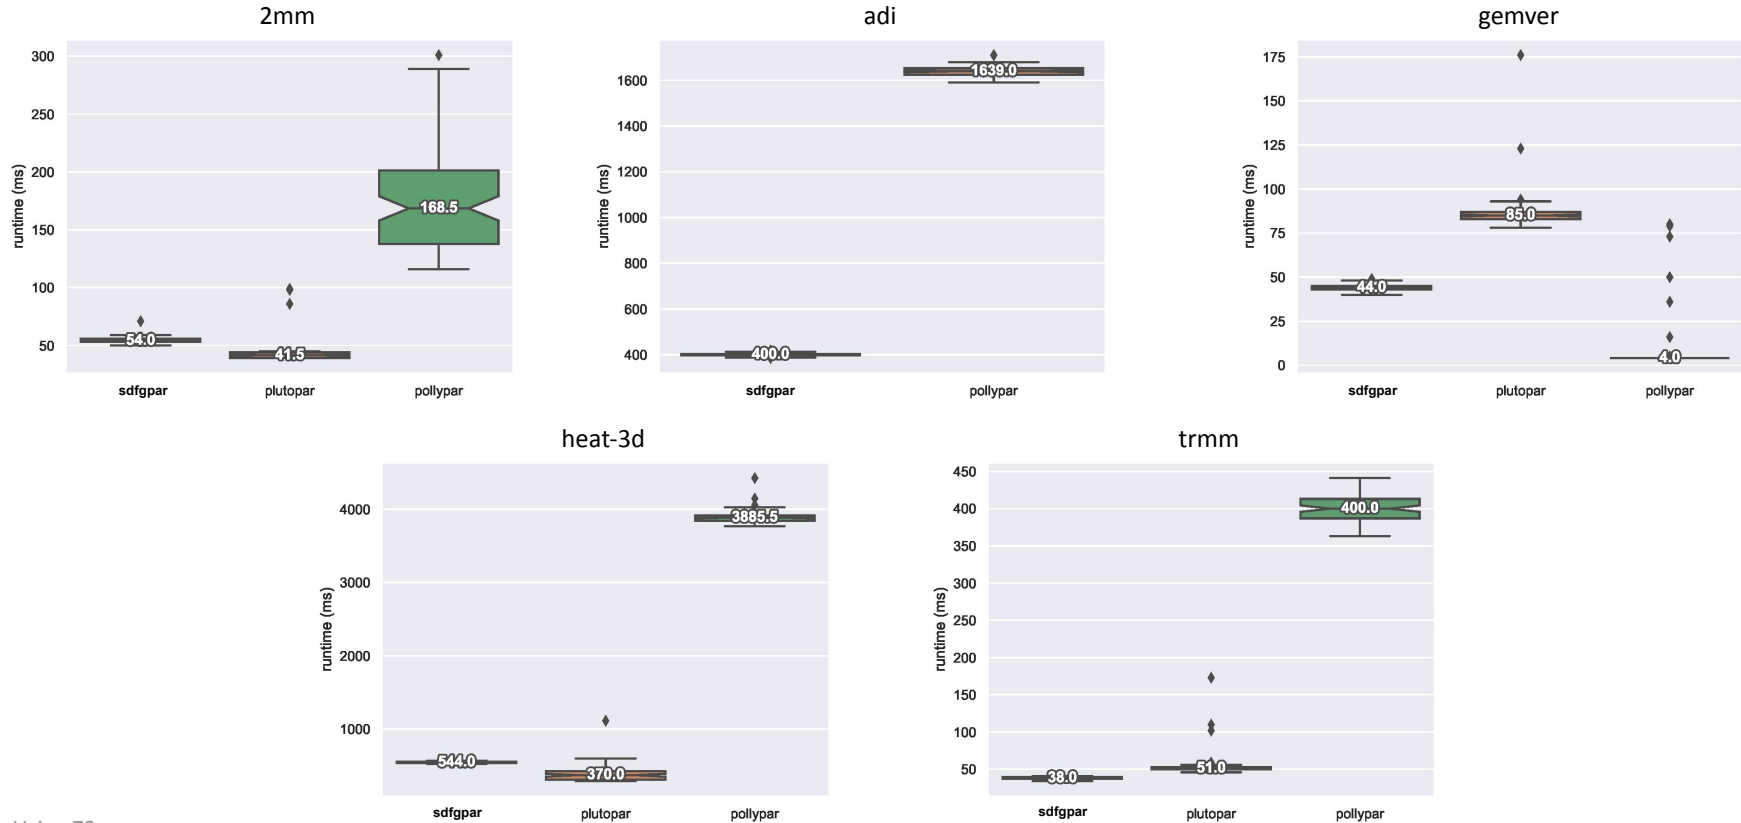
2MM Affine vs. MLIR-SDFG



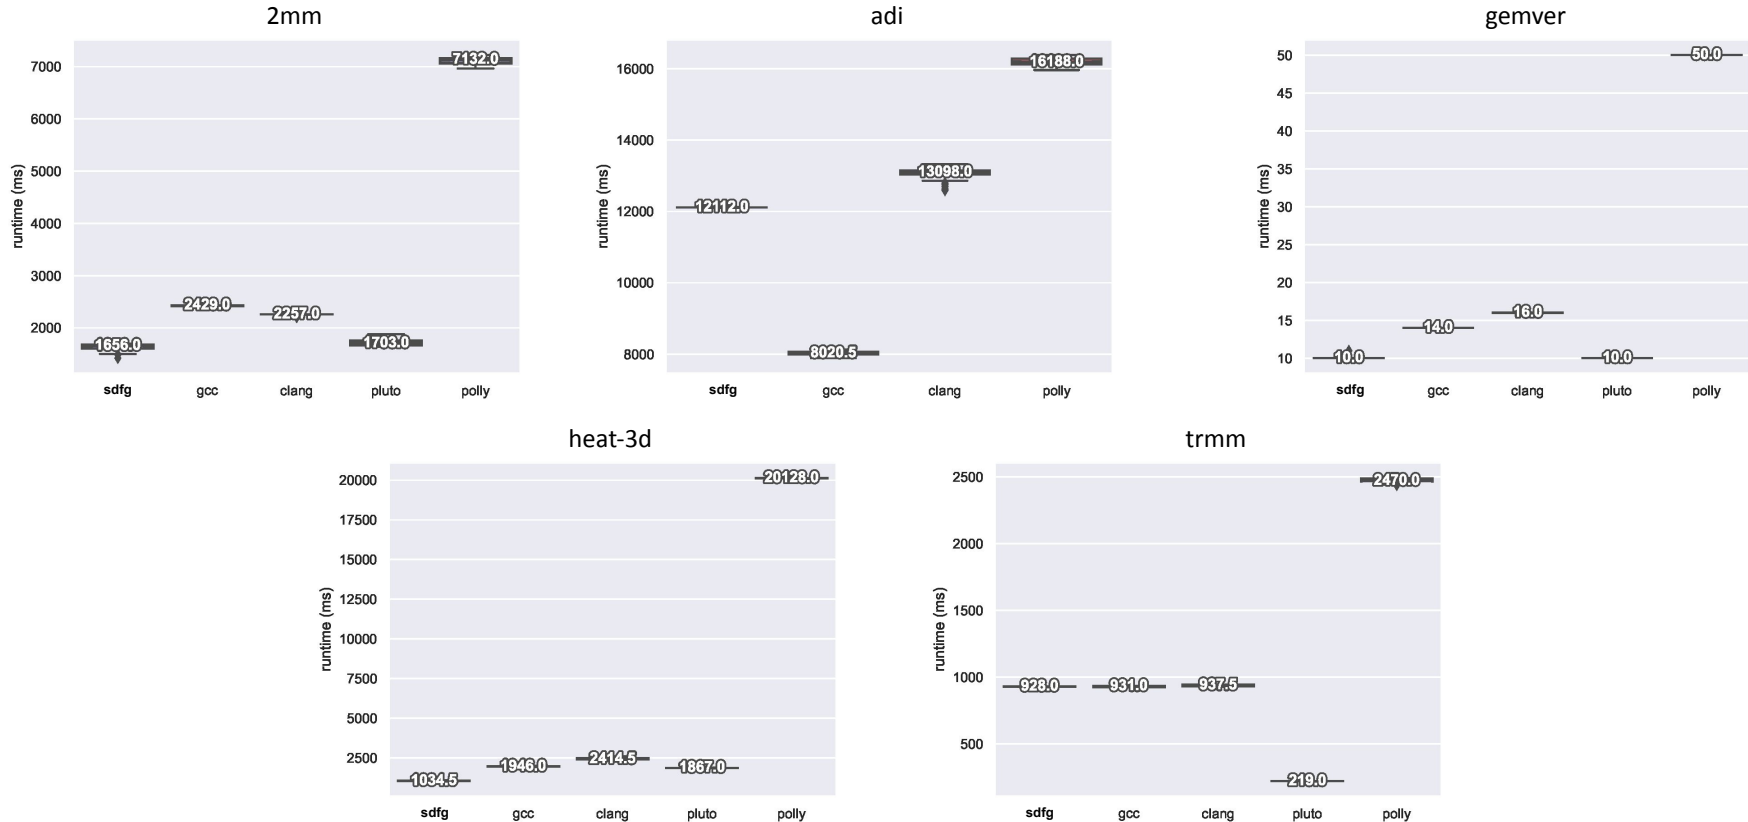
2MM Affine vs. MLIR-SDFG



Parallel Benchmarks




Sequential Benchmarks



Open MLIR Challenges

```
func private @kernel_2mm(%arg0: i32, ..., %arg6: memref<?x900xf64>, ...){  
  ...  
  %5 = arith.index_cast %arg0 : i32 to index  
  ...  
  scf.for %arg11 = %c0_0 to %5 step %c1_1 {  
    %8 = memref.load %arg6[%arg11, %arg13] : memref<?x900xf64>  
  }  
  ...  
}
```

Open MLIR Challenges



```
func private @kernel_2mm(%arg0: i32, ..., %arg6: memref<?x900xf64>, ...){  
  ...  
  %5 = arith.index_cast %arg0 : i32 to index  
  ...  
  scf.for %arg11 = %c0_0 to %5 step %c1_1 {  
    %8 = memref.load %arg6[%arg11, %arg13] : memref<?x900xf64>  
  }  
  ...  
}
```

Open MLIR Challenges

```
func private @kernel_2mm(%arg0: i32, ..., %arg6: memref<?x900xf64>, ...)
```



```
!sdfg.array<sym("s_0")x900xf64>
```

Open MLIR Challenges

```
func private @kernel_2mm(%arg0: i32, ..., %arg6: memref<?x900xf64>, ...)
```




```
!sdfg.array<sym("s_0")x900xf64>
```



What's the value?

Open MLIR Challenges

```
func private @kernel_2mm(%arg0: i32, ..., %arg6: memref<?x900xf64>, ...)
```

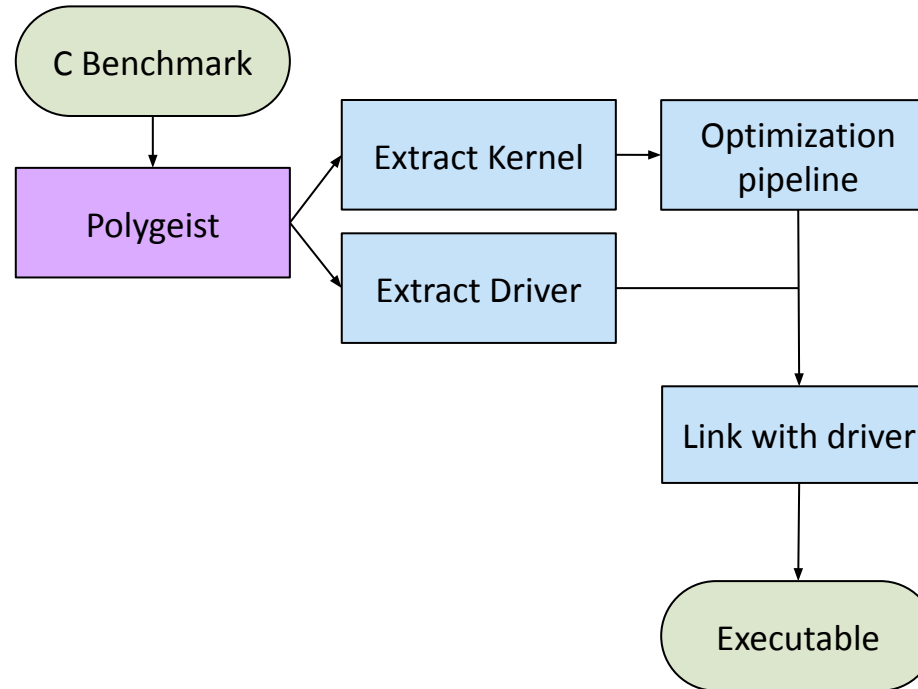


```
!sdfg.array<sym("s_0")x900xf64>
```

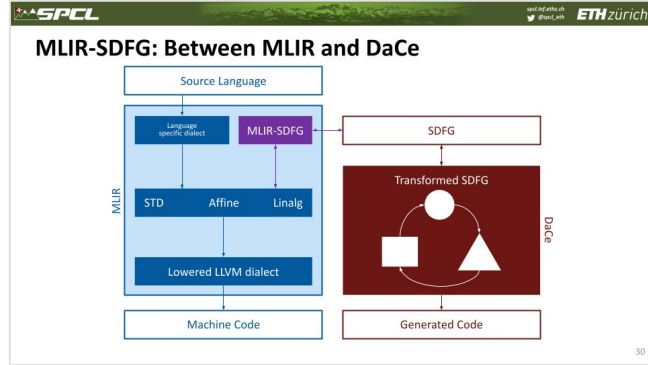


What's the value?

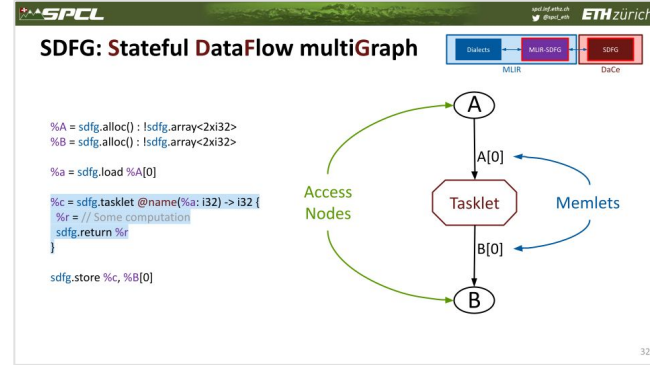
Future Work



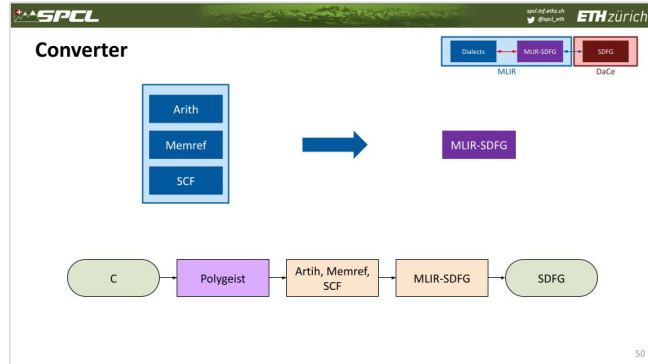
Summary




30



32



50



Using 72 cores

72

Berke Ates, Prof. Dr. Torsten Hoefler, Dr. Tal Ben-Nun, Dr. Alexandru Calotoiu

MLIR-SDFG: A Data-Centric Dialect for MLIR

